

Group 3

Joshua Lopez, Alex Castillo, Edward Cruz, Gerardo Gonzalez

REPORT

Executive Summary

This business case outlines the development of an NLP model to automate the processing of customer feedback for a retail company. The goal is to classify customer reviews into positive, negative, or neutral categories to help the company improve its products and services. The second part is to use GenerativeAI to summarize reviews broken down into review score (0-5), and broken down into product categories - if the categories are too many to handle, select a top-K categories. Create a clickable and dynamic visualization dashboard using a tool like Tableau, Plotly, or any of your choice.

Project goals

- The ML/AI system should be able to run classification of customers' reviews (the textual content of the reviews) into positive, neutral, or negative.
- For a product category, create a summary of all reviews broken down by each star or rating (we should have 5 of these).
 - If your system can't handle all products categories, pick a number that you can work with (eg top 10, top 50, Etc)

1) Traditional NLP & ML approaches


Dateset used: 1429_1.csv

From:<https://www.kaggle.com/datasets/datafiniti/consumer-reviews-of-ama-zon-products>

a) Data Preprocessing:

- We analyzed the dataset to determine which columns were really useful for the case.
- Check features that had null values and clean them.
- After removing columns tat had most null values. By checking the distributions of the other features we determined that for the sentiment classifier, the most important feature are “reviews.text” (model input features) and “reviews.rating” (model target) .

Noted: From the analysis we found that the dataset is very unbalance.




	reviews.rating	No of Users
0	5.0	23775
1	4.0	8541
2	3.0	1499
3	1.0	410
4	2.0	402

b) Clean Texts Tokenize and Remove stopwords

- i) First split the dataset into features and target
- ii) To the input feature column (review.text):
 - Clean stop words (NLTK)
 - Convert to lowercase
 - Remove punctuations
- iii) To the target column (reviews.rating):
 - creates a function to map rating to sentiments.
 - Positive (ratings 4-5)
 - Negative (ratings 1-2)
 - Neutral (ratings 3)

c) Vectorize text data using TF-IDF and Balancing

- Vectorize the clean review.text.
- I apply SMOTE on the vectorize data to balance the dataset.



d) Split data into Training and Test and Train using Different Traditional NLP and ML models

1) Naive Bayes

Accuracy: 0.84

Classification Report:

	precision	recall	f1-score	support
negative	0.86	0.89	0.87	6510
neutral	0.80	0.77	0.79	6388
positive	0.86	0.86	0.86	6488
accuracy			0.84	19386
macro avg	0.84	0.84	0.84	19386
weighted avg	0.84	0.84	0.84	19386

Confusion Matrix:

```
[[5775  495  240]
 [ 793 4916  679]
 [ 179  701 5608]]
```

2) Random Forest


Parameter:

```
# Create and train the Random Forest model
rf_classifier = RandomForestClassifier(
    #class_weight='balanced',
    criterion='gini',
    n_estimators=100, # number of trees
    max_depth=None,  # maximum depth of trees
    min_samples_split=2,
    min_samples_leaf=1,
    random_state=42,
    warm_start = False,
    verbose = 1     # for reproducibility
)
```

Accuracy: 0.99

Random Forest Classification Report:

	precision	recall	f1-score	support
negative	0.99	0.99	0.99	6510
neutral	0.98	0.99	0.98	6388
positive	0.99	0.98	0.99	6488
accuracy			0.99	19386
macro avg	0.99	0.99	0.99	19386
weighted avg	0.99	0.99	0.99	19386



```
Cross-validation scores: [0.98084983 0.98084859 0.98091308 0.98284756 0.9800748 ]
Average CV score: 0.981106772635683
CV score standard deviation: 0.0009236352971816066
```

Analysis

These cross-validation results:

1. Consistency across folds:

- The scores across all 5 folds are very consistent (ranging from 0.980 to 0.982)
- The standard deviation is very small (0.00092), which is excellent
- This consistency suggests that your model is stable and performs similarly across different subsets of the data

2. Average CV Score:

- Average score of 0.981 (98.1%) is slightly lower than your test set accuracy of 0.99
- This small difference (about 0.9%) between CV and test performance suggests that your model isn't severely overfitting
- It's performing consistently well across different data splits

3. Analysis:

- The high scores with low variance suggest your model is genuinely learning the patterns in your data
- The small gap between CV scores and test accuracy indicates minimal overfitting
- The extremely low standard deviation (0.00092) shows remarkable stability in model performance

Conclusion:

Based on these cross-validation results, I would say your model is NOT overfitting. The reasons are:

1. Very stable performance across different data splits

2. Small difference between CV and test performance
3. Extremely low standard deviation in CV scores

Your Random Forest model appears to be genuinely good at this classification task, likely because:

- The relationship between review text and sentiment might be relatively straightforward
- You probably have a good feature representation
- The classes are well-balanced after your preprocessing

3) Random Forest L2 Regularizer


Parameter: which parameter is use????

```
rf_regularized = RandomForestClassifier(  
    n_estimators=100,  
    max_depth=None,           # Limit tree depth  
    min_samples_split=5,      # Require more samples to split  
    min_samples_leaf=4,       # Require more samples in leaves  
    max_features='sqrt',      # Reduce features considered per split  
    max_samples=0.8,          # Use bootstrapping with 80% of samples  
    random_state=42  
)
```

Report:

```
Accuracy: 0.91  
  
Random Forest Classification Report:  
              precision    recall  f1-score   support  
  
   negative      0.94      0.95      0.95     6510  
    neutral      0.92      0.87      0.89     6388  
   positive      0.88      0.93      0.91     6488  
  
   accuracy                    0.91     19386  
  macro avg      0.92      0.91      0.91     19386  
 weighted avg      0.92      0.91      0.91     19386
```

Confusion Matrix:



Cross-Validation:


```
Cross-validation scores: [0.89941324 0.90346918 0.90063193 0.9048878 0.89760124]
Average CV score: 0.9012006781887552
CV score standard deviation: 0.0026538770363785692
```


4) GradientBoosting

Report:

Accuracy: 0.66				
Gradient Boosting Classification Report:				
	precision	recall	f1-score	support
negative	0.77	0.68	0.72	6510
neutral	0.52	0.71	0.60	6388
positive	0.78	0.61	0.68	6488
accuracy			0.66	19386
macro avg	0.69	0.66	0.67	19386
weighted avg	0.69	0.66	0.67	19386

Confusion Matriz:




5) XgBOOST

Report:

XGBoost Classification Report:				
	precision	recall	f1-score	support
negative	0.87	0.83	0.85	6510
neutral	0.80	0.74	0.77	6388
positive	0.78	0.88	0.83	6488
accuracy			0.82	19386
macro avg	0.82	0.82	0.81	19386
weighted avg	0.82	0.82	0.82	19386

Confusion Matrix:

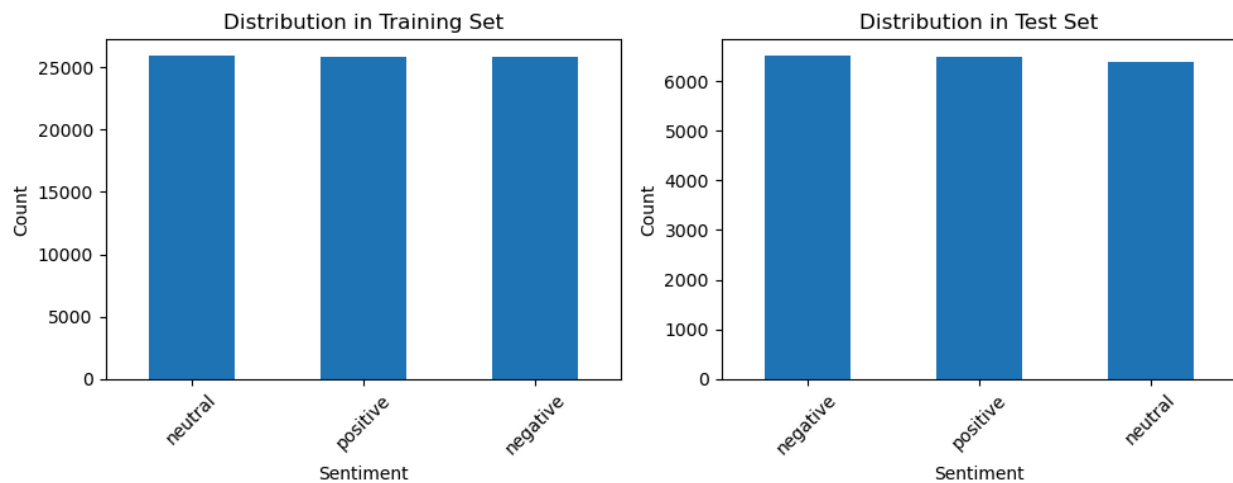


2) Sequence-to-Sequence modeling with LSTM

- **Goal:** Build a Bidirectional LSTM model to predict the review class i.e., negative, positive, or neutral.

a) I did the same data preprocessing as for the traditional NPL models:

- Deleted unnecessary columns
- Remove nulls
- Remove stopwords, puntuacions.
- Convert to lowercase
- Word tokenize
- Vectorize the data
- Balance the data with SMOTE
- Split data into training and validation datasets



b) Train LSTM MODEL

- Reshape input data for LSTM (samples, timesteps, features)
- Architecture:

```
# Create LSTM Model
model = Sequential([
    LSTM(64, input_shape=(timesteps, input_dim), return_sequences=False),
    Dropout(0.5),
    Dense(32, activation='relu'),
    Dropout(0.5),
    Dense(3, activation='softmax') # 3 classes: negative, neutral, positive
])
```

```
model.compile(
    optimizer=tf.keras.optimizers.Adam(learning_rate=0.001),
    loss='sparse_categorical_crossentropy',
    metrics=['accuracy']
)
```

```
# Add early stopping
early_stopping = EarlyStopping(
    monitor='val_loss',
    patience=5,
    restore_best_weights=True
)
```

```
# Train the model
history = model.fit(
    X_train_resaped,
    y_train,
    epochs=20,
    batch_size=64,
    validation_split=0.2,
    callbacks=[early_stopping]
)
```

c) Evaluation

-Report:


```
Overall Accuracy: 94.82%

Classification Report:



```

	precision	recall	f1-score	support
Negative	0.95	1.00	0.97	6462
Neutral	0.92	0.96	0.94	6462
Positive	0.98	0.88	0.93	6462
accuracy			0.95	19386
macro avg	0.95	0.95	0.95	19386
weighted avg	0.95	0.95	0.95	19386

-Confision Matrix:



-Accuracy and Loss plots



3) Transformer approach (HuggingFace API)

The goal was to have a **summary of all reviews broken down by each star or rating, for all product categories**. It would look like this, where every product had, for each rating, a single summarization of all reviews combined of that rating.

Product Category	Stars	Summaries of reviews
Electronics	★	Summary of all 1 star reviews
	★ ★	Summary of all 2 star reviews
	★ ★ ★	Summary of all 3 star reviews
	★ ★ ★ ★	Summary of all 4 star reviews
	★ ★ ★ ★ ★	Summary of all 5 star reviews
Electronics,Media	★	Summary of all 1 star reviews
	★ ★	Summary of all 2 star reviews
	★ ★ ★	Summary of all 3 star reviews

The dataset used:

Datafiniti_Amazon_Consumer_Reviews_of_Amazon_Products.csv

from <https://www.kaggle.com/datasets/datafiniti/consumer-reviews-of-amazon-products>

```
kaggle_df = pd.read_csv('Datafiniti_Amazon_Consumer_Reviews_of_Amazon_Products.csv') # Load Dataset
```

The model used: **t5-small** from <https://huggingface.co/google-t5/t5-small>

```
# Initialize the model and tokenizer
model_name = "t5-small" # model
model = T5ForConditionalGeneration.from_pretrained(model_name, device_map={"": 0})

tokenizer = T5Tokenizer.from_pretrained(model_name)
```

Data Preprocessing

1.1 Data Cleaning and Tokenization

Cleaning NULLS

Since the columns that were going to be used were “**categories**”, “**review.rating**” and “**review.text**”, and they didn’t have **NULL** values, eliminating the **NULL** fields was not needed.

Clean Nulls

```
print(kaggle_df.isnull().sum()) # No NULLs found in categories, reviews.text or reviews.rating  
✓ 0.0s
```

id	0
dateAdded	0
dateUpdated	0
name	0
asins	0
brand	0
categories	0
primaryCategories	0
imageURLs	0
keys	0
manufacturer	0
manufacturerNumber	0
reviews.date	0
reviews.dateAdded	3948
reviews.dateSeen	0
reviews.doRecommend	0
reviews.id	4971
reviews.numHelpful	0
reviews.rating	0
reviews.sourceURLs	0
reviews.text	0
reviews.title	13
reviews.username	1
sourceURLs	0

dtype: int64

Tokenization

For the tokenization the T5 library provided an instruction to encode the data to be summarized.

```
tokenizer = T5Tokenizer.from_pretrained(model_name)

# Join reviews into a single string
text = "\n\n".join(dataframe)

# Tokenize and summarize the input text. inputs is a pytorch tensor, torch.Tensor
inputs = tokenizer.encode("summarize: " + text, return_tensors = "pt", truncation = True).to("cuda:0")
```

Metrics

Summarization models use **ROUGE metrics** instead of accuracy scores to verify how good the model is. The ROUGE scores used were **rouge1**, **rouge2**, **rougeL** and **rougeLsum**. It is preferred to use **rougeL** since it uses the **longest common sequence** to produce its score. Since its **rougeL** score is 0.087, it is not a good model to use.

The **scores** for the **pre-trained** model are:

rouge1 average: 0.0873561269402484
rouge2 average: 0.0
rougeL average: 0.08728537224163033
rougeLsum average: 0.0873561269402484

```
{'rouge1': 0.05319148936170213, 'rouge2': 0.0, 'rougeL': 0.05319148936170213, 'rougeLsum': 0.05319148936170213}
{'rouge1': 0.0390625, 'rouge2': 0.0, 'rougeL': 0.0390625, 'rougeLsum': 0.0390625}
rouge1 average: 0.0873561269402484 - rouge2 average: 0.0 - rougeL average: 0.08728537224163033 - rougeLsum average: 0.0873561269402484
```

End Result:

```
category_dict_df = pd.DataFrame(columns = ['Product Category', 'Rating', 'Summary of reviews'])

for key in category_dict: # for each product category
    for rating in range(1, 6): # from 1 stars to 5 stars, rating
        category_dict_df.loc[len(category_dict_df)] = [key, rating, category_dict[key][rating - 1]]

display(category_dict_df)
```

✓ 0.0s

	Product Category	Rating	Summary of reviews
0	Computers,Electronics Features,Tablets,Electro...	1	the whitepaper looks Identical to the \$120 mod...
1	Computers,Electronics Features,Tablets,Electro...	2	screen too dark The screen is too dark, and ca...
2	Computers,Electronics Features,Tablets,Electro...	3	NULL
3	Computers,Electronics Features,Tablets,Electro...	4	the kindle is good to download apps for books ...
4	Computers,Electronics Features,Tablets,Electro...	5	the amazon Kindle is light weight and easy to ...
...
110	Tablets,Fire Tablets,Electronics,iPad & Tablet...	1	very cheap and was not impressed at all never ...
111	Tablets,Fire Tablets,Electronics,iPad & Tablet...	2	NULL
112	Tablets,Fire Tablets,Electronics,iPad & Tablet...	3	the battery is having more and more trouble ho...
113	Tablets,Fire Tablets,Electronics,iPad & Tablet...	4	my daughter has had this tablet for almost 2 m...
114	Tablets,Fire Tablets,Electronics,iPad & Tablet...	5	NULL

115 rows × 3 columns

Fine Tuning

The dataset used: 'gopalkalpande/bbc-news-summary'
from **Dataset** library

The model used: **t5-small** from <https://huggingface.co/google-t5/t5-small>

Fine-Tuning

```
dataset = load_dataset('gopalkalpande/bbc-news-summary', split = 'train')
full_dataset = dataset.train_test_split(test_size = 0.2, shuffle = True)

dataset_train = full_dataset['train'] # full_dataset['train'] # text?
dataset_valid = full_dataset['test'] # cambiar por category_dict[all_cate

print(dataset_train)
print(dataset_valid)
```

✓ 4.1s

```
Dataset({
  features: ['File_path', 'Articles', 'Summaries'],
  num_rows: 1779
})
Dataset({
  features: ['File_path', 'Articles', 'Summaries'],
  num_rows: 445
})
```

Parameters

```
MODEL = 't5-small'  
BATCH_SIZE = 4  
NUM_PROCS = 4  
EPOCHS = 10  
OUT_DIR = 'results_t5small'  
MAX_LENGTH = 512 # Maximum c
```

✓ 0.0s

```
training_args = TrainingArguments(  
    output_dir = OUT_DIR,  
    num_train_epochs = EPOCHS, # number of epochs  
  
    per_device_train_batch_size = BATCH_SIZE,  
    per_device_eval_batch_size = BATCH_SIZE,  
  
    warmup_steps = 500,  
  
    weight_decay = 0.01,  
  
    evaluation_strategy = 'steps', # how often will evaluation be done  
    eval_steps = 200,  
  
    save_strategy = 'epoch', # how often will saving be during training  
    save_total_limit = 2,  
  
    learning_rate = 0.001,  
    # dataloader_num_workers = 4 # Number of subprocesses to use for data loading  
)
```

```
trainer = Trainer(  
    model = model,  
  
    args = training_args,  
  
    train_dataset = tokenized_train,  
    eval_dataset = tokenized_valid,  
  
    preprocess_logits_for_metrics = preprocess_logits_for_metrics,  
    compute_metrics = compute_metrics # The function that will be used to compute the metrics  
)
```

Metrics

Rouge1	Rouge2	RougeL
0.898500	0.828500	0.881800

The **scores** for the **fine-tuned** model were:

rouge1: 0.898500

rouge2: 0.828500

rougeL: 0.881800

It had **10x** a better **rougeL score** than the previous model. So this model would be a whole lot more usable for text summarization.

Transformer Approach for Sentiment Classification:

1) Foundation model (No Finetuning)

-For this case, we use the BERT model. This model has already been pre-trained to classify product reviews into ratings.

Dataset: 1429_1.csv

from <https://www.kaggle.com/datasets/datafiniti/consumer-reviews-of-amazon-products>

Foundation Model: bert-base-uncased

Data Preprocessing:

- Clean nulls
- Drop unnecessary columns
- Remove short and super long reviews.
- Group together rating and Map them to Sentiment: Rating 1-2 = Negative,

Rating 3 = neutral and rating 4-5 = positive.


Load Model and Tokenizer

```
# Load models
print("Loading BERT sentiment model...")
tokenizer_sentiment = AutoTokenizer.from_pretrained('bert-base-uncased')
model_sentiment = AutoModelForSequenceClassification.from_pretrained('bert-base-uncased')
```

Classification Report:

Classification Report:				
	precision	recall	f1-score	support
negative	0.395	0.755	0.518	730
neutral	0.288	0.419	0.341	1404
positive	0.980	0.939	0.959	31372
accuracy			0.913	33506
macro avg	0.554	0.704	0.606	33506
weighted avg	0.938	0.913	0.924	33506

Confusion Matrix:



2) Fine tuned Model (BERT)

-We decided to do full fine-tune after not improving the result with the LoRA Configuration. Since our dataset has 33,506 reviews and we have the computational resources.

Data Preprocessing:

-Similar data preprocess as in the Foundation model.

How data distribution looks before grouping them into sentiment:

```
Distribución de ratings:
reviews.rating
1.0      367
2.0      363
3.0     1404
4.0     8200
5.0    23172
Name: count, dtype: int64

Distribución de sentimientos:
sentiment
0         730
1        1404
2       31372
Name: count, dtype: int64
```

Split data into Training and Validation

```
# Dividir los datos
train_texts, val_texts, train_labels, val_labels = train_test_split(
    df['reviews.text'].tolist(),
    df['sentiment'].tolist(),
    test_size=0.2,
    random_state=42
)
```

Load Model and Tokenizer:


```
# Cargar el tokenizer y modelo BERT
tokenizer = AutoTokenizer.from_pretrained('bert-base-uncased')
model = AutoModelForSequenceClassification.from_pretrained(
    'bert-base-uncased',
    num_labels=3
)
```

Configure dataloader and Train model:


```
# Configurar dataloaders con un tamaño de batch más pequeño
train_loader = DataLoader(train_dataset, batch_size=32, shuffle=True)
val_loader = DataLoader(val_dataset, batch_size=32)

# Entrenamiento
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
model.to(device)
optimizer = torch.optim.AdamW(model.parameters(), lr=2e-5)
```

Evaluation using the full dataset:



```
=====
Full Dataset Evaluation Results:
=====
Average Loss: 0.0726
Accuracy: 0.9835
F1 Score (weighted): 0.9831

Detailed Classification Report:
-----

```

	precision	recall	f1-score	support
0	0.90	0.88	0.89	730
1	0.86	0.78	0.82	1404
2	0.99	1.00	0.99	31372
accuracy			0.98	33506
macro avg	0.92	0.88	0.90	33506
weighted avg	0.98	0.98	0.98	33506